

CS137 FUSE Documentation

There is very little FUSE documentation on the [FUSE Web site](#). A bit more, which is unfortunately very outdated, is available from [an IBM article from 2006](#). If you come across anything more complete or more current, I'd appreciate hearing about it so I can add a link to it from this site.

Note: Be sure to read the [Gotchas](#) list before starting your code, and refer back to it frequently when you run into troubles.

Writing a FUSE Client

The best way to write a fuse client is to start with an example or an existing client; I recommend [fusexmp.c](#) or [fusexmp_fh.c](#) (the latter implements file handles, so it's a better choice if you're developing a complex filesystem). The existing clients provide a scaffolding for you to work from, but you'll still need to understand [what all the functions are supposed to do](#), and also how to [compile](#) and [run](#) your client. That's what this Web page is for.

Unix Manual Pages

Many of the FUSE functions are closely related to Unix system calls. Rather than repeating the full specification (especially error conditions) here, it's better for you to refer to the Unix man page. You can do this on any Unix machine with the "man" command. By convention, if I refer you to the "stat(2) system call", that means you should type "man 2 stat" to get the necessary information.

FUSE File Handles

Many FUSE functions offer two ways to identify the file being operated upon. The first, which is always available, is the "path" argument, which is the full pathname (relative to the filesystem root) of the file in question. If you choose to do so, all your functions can use that argument to locate a file.

However, pathname lookup is often a very expensive operation, so FUSE sometimes provides you another option: a "file handle" in the "fuse_file_info" structure. The file handle is stored in that structure's "fh" element, which is an unsigned 64-bit integer (`uint64_t`) interpreted by FUSE. If you choose to use it, you should set that field in your [open](#) and [opendir](#) functions; other functions can then use it. In many FUSE implementations, the file handle is actually a pointer to a useful data structure, which is typecast to an integer only to keep the compiler happy. But you can make it an index into an array, a hash key, or pretty much anything else you choose.

Getting FUSE Context

For many operations, it is useful to have a relevant "context" in which to execute them. For historical reasons, the context isn't passed as an argument; instead you must call `fuse_get_context` with no argument, which returns a pointer to a `struct fuse_context` with the following usable elements:

```
uid          The (numeric) user ID of the process invoking the operation.

gid          The (numeric) group ID of the process invoking the operation.

pid         The thread (process) ID of the process invoking the operation.

private_data A pointer (void*) to the private data returned by the init function.

umask       The umask of the process invoking the operation. See umask\(2\).
```

FUSE Functions

The following is a brief description of all the API functions you can create in a FUSE filesystem. Note that many are not required, especially if you are implementing a partial filesystem like the ones in these assignments; especially important functions are highlighted in yellow. However, I have tried to provide documentation for all the functions here. Unless otherwise specified, all functions return an integer 0 or a positive number for success, or a negative value selected from `errno.h` if there is an error.

All of your functions should be named with a prefix that is closely related to your filesystem name. For example, in an SSH filesystem you should use `ssh_getattr`, `ssh_read`, etc.

```
void* init(struct fuse_conn_info *conn)
    Initialize the filesystem. This function can often be left unimplemented, but it can be a handy way to perform one-time setup such as allocating variable-sized data structures in memory or initializing the persistent storage (aka "disks"). The fuse_conn_info structure gives information about what features are supported by FUSE, and can be used to request certain capabilities (see "Init Function" below for more information). The return value from init will be made available to all file operations in the private_data field of fuse_context. It is also passed as a parameter to the destroy() method. (Note: see the warning under Other Options below, regarding relative pathnames.)

void destroy(void* private_data)
    Called when the filesystem exists. The private_data comes from the return value of init. This function should free any data that init allocated with malloc, and should make sure everything related to the filesystem has been flushed to persistent storage.

getattr(const char* path, struct stat* stbuf)
    Return file attributes. The "stat" structure is described in detail in the stat(2) manual page. For the given pathname, this should fill in the elements of the "stat" structure. If a field is meaningless or semi-meaningless (e.g., st_rdev) then it should be set to 0 or given a "reasonable" value. This call is pretty much required for a usable filesystem. NOTE: st_dev and st_ino are often used by application programs to decide whether two file names are aliases for the same physical storage, so setting them to 0 isn't ideal.

fgetattr(const char* path, struct stat* stbuf, struct fuse_file_info *fi)
    As getattr, but called when fgetattr(2) is invoked by the user program. Usually it can just call getattr.

access(const char* path, mask)
    This is the same as the access(2) system call. It returns -EENOENT if the path doesn't exist, -EACCESS if the requested permission isn't available, or 0 for success. Note that it can be called on files, directories, or any other object that appears in the filesystem. This call is not required but is strongly recommended.

readlink(const char* path, char* buf, size_t size)
    If path is a symbolic link, fill buf with its target, up to size. See readlink(2) for how to handle a too-small buffer and for error codes. Not required if you don't support symbolic links. NOTE: Symbolic-link support requires only readlink and symlink. FUSE itself will take care of tracking symbolic links in paths, so your path-evaluation code doesn't need to worry about it.

opendir(const char* path, struct fuse_file_info* fi)
    Open a directory for reading.

readdir(const char* path, void* buf, fuse_fill_dir_t filler, off_t offset, struct fuse_file_info* fi)
    Return one or more directory entries (struct dirent) to the caller. This is one of the most complex FUSE functions. It is related to, but not identical to, the readdir(2) and getdents(2) system calls, and the readdir(3) library function. Because of its complexity, it is described separately in "Readir Function" below. Required for essentially any filesystem, since it's what makes ls and a whole bunch of other things work.

mknod(const char* path, mode_t mode, dev_t rdev)
    Make a plain file, special (device) file, FIFO, or socket. See mknod(2) for most details. This function is called to create new files; if mode satisfies S_IRRSG then mknod should enter an empty file into its parent directory. For other file types, this function is rarely needed, since it's uncommon to make these objects inside special-purpose filesystems. In my implementation, if S_IRRSG(mode) is false then I just return -EACCESS to reject the call.

create(const char* path, mode_t mode)
    Make a plain file. This is a less flexible version of mknod. I implemented it in one line by passing the call off to mknod and adding S_IFREG to the mode.

mkdir(const char* path, mode_t mode)
    Create a directory with the given name. The directory permissions are encoded in mode. See mkdir(2) for details. This function is needed for any reasonable read/write filesystem.

unlink(const char* path)
    Remove (delete) the given file, symbolic link, hard link, or special node—but NOT a directory. Note that if you support hard links, unlink only deletes the data when the last hard link is removed. See unlink(2) for details. This function is needed for almost any read/write filesystem.

rmdir(const char* path)
    Remove the given directory. This should succeed only if the directory is empty (except for "." and ".."). See rmdir(2) for details.

symlink(const char* path, const char* from)
    Create a symbolic link named "from" which, when evaluated, will lead to "to". Not required if you don't support symbolic links. NOTE: Symbolic-link support requires only readlink and symlink. FUSE itself will take care of tracking symbolic links in paths, so your path-evaluation code doesn't need to worry about it.

rename(const char* from, const char* to)
    Rename the file, directory, or other object "from" to the target "to". Note that the source and target don't have to be in the same directory, so it may be necessary to move the source to an entirely new directory. Also note that if the target already exists, it must be atomically replaced; many application programs depend on this behavior. See rename(2) for full details.

link(const char* from, const char* to)
    Create a hard link between "from" and "to". Hard links aren't required for a working filesystem, and many successful filesystems don't support them. If you do implement hard links, be aware that they have an effect on how unlink works. See link(2) for details.

chmod(const char* path, mode_t mode)
    Change the mode (permissions) of the given object to the given new permissions. Only the permissions bits of mode should be examined. See chmod(2) for details.

chown(const char* path, uid_t uid, gid_t gid)
    Change the given object's owner and group to the provided values. See chown(2) for details. NOTE: FUSE doesn't deal particularly well with file ownership, since it usually runs as an unprivileged user and this call is restricted to the superuser. It's often easier to pretend that all files are owned by the user who mounted the filesystem, and to skip implementing this function.

truncate(const char* path, off_t size)
    Truncate or extend the given file so that it is precisely size bytes long. See truncate(2) for details. This call is required for read/write filesystems, because recreating a file will truncate it.

ftruncate(const char* path, off_t size)
    As truncate, but called when ftruncate(2) is called by the user program. Usually this can just call truncate.

utimens(const char* path, const struct timespec ts[2])
    Update the last access time of the given object from ts[0] and the last modification time from ts[1]. For full time specifications are given to nanosecond resolution, but your filesystem doesn't have to be that precise; see utimensat(2) for full details. Note that the time specifications are allowed to have certain special values; however, I don't know if FUSE functions have to support them. This function isn't necessary but is nice to have in a fully functional filesystem.

open(const char* path, struct fuse_file_info* fi)
    Open a file. If you aren't using file handles, this function should just check for existence and permissions and return either success or an error code. If you use file handles, you should also allocate any necessary structures and set fi->fh. In addition, fi has some other fields that an advanced filesystem might find useful; see the structure definition in fuse_common.h for very brief commentary.

read(const char* path, char* buf, size_t size, off_t offset, struct fuse_file_info* fi)
    Read size bytes from the given file into the buffer buf, beginning offset bytes into the file. See read(2) for full details. Returns the number of bytes transferred, or 0 if offset was at or beyond the end of the file. Required for any sensible filesystem.

write(const char* path, const char* buf, size_t size, off_t offset, struct fuse_file_info* fi)
    As for read above, except that it can't return 0.

stats(const char* path, struct statvfs* stbuf)
    Return statistics about the filesystem. See statvfs(2) for a description of the structure contents. Usually, you can ignore the path. Not required, but handy for read/write filesystems since this is how programs like df determine the free space.

release(const char* path, struct fuse_file_info* fi)
    This is the only FUSE function that doesn't have a directly corresponding system call, although close(2) is related. release is called when FUSE is explicitly done with a file; at that point, you can free up any temporarily allocated data structures. The IBM document claims that there is exactly one release per open, but I don't know if that is true.

releasedir(const char* path, struct fuse_file_info* fi)
    This is like release, except for directories.

fsync(const char* path, int isdatasync, struct fuse_file_info* fi)
    Flush any dirty information about the file to disk. If isdatasync is nonzero, only data, not metadata, needs to be flushed. When this call returns, all file data should be on stable storage. Many filesystems leave this call unimplemented, although technically that's a Bad Thing since it risks losing data. If you store your filesystem inside a plain file on another filesystem, you can implement fsync by calling fdatsync(2) on that file, which will flush too much data (slowing performance) but achieve the desired guarantee. A higher-performance option would be to open the backing file with O_DSYNC and manage a private cache, but this approach is too complex for most FUSE filesystems.

fsyncdir(const char* path, int isdatasync, struct fuse_file_info* fi)
    Like fsync, but for directories.

flush(const char* path, struct fuse_file_info* fi)
    Called on each close so that the filesystem has a chance to report delayed errors. Important: there may be more than one flush call for each open. Note: There is no guarantee that flush will ever be called at all!

lock(const char* path, struct fuse_file_info* fi, int cmd, struct flock* locks)
    Perform a POSIX file-locking operation. See "Lock Function" below.

bmap(const char* path, size_t blocksize, uint64_t* blockno)
    This function is similar to bmap(9). If the filesystem is backed by a block device, it converts blockno from a file-relative block number to a device-relative block number. It isn't entirely clear how the blocksize parameter is intended to be used.

setxattr(const char* path, const char* name, const char* value, size_t size, int flags)
    Set an extended attribute. See setxattr(2). This function should be implemented only if the preprocessor constant HAVE_SETXATTR is defined.

getxattr(const char* path, const char* name, char* value, size_t size)
    Read an extended attribute. See getxattr(2). This function should be implemented only if the preprocessor constant HAVE_SETXATTR defined.

listxattr(const char* path, const char* list, size_t size)
    List the names of all extended attributes. See listxattr(2). This function should be implemented only if the preprocessor constant HAVE_SETXATTR is defined.

ioctl(const char* path, int cmd, void* arg, struct fuse_file_info* fi, unsigned int flags, void* data)
    Support the ioctl(2) system call. As such, almost everything is up to the filesystem. On a 64-bit machine, FUSE_IOCTL_COMPAT will be set for 32-bit ioctls. The size and direction of data being passed by _IOC_*() decoding of cmd. For _IOC_NONE, data will be NULL; for _IOC_WRITE data is being written by the user, for _IOC_READ it is being read, and if both are set the data is read/written in cases where the data area is _IOC_SIZE(cmd) bytes in size.

poll(const char* path, struct fuse_file_info* fi, struct fuse_pollhandle* ph, unsigned* revents)
    Poll for I/O readiness. If ph is non-NULL, when the filesystem is ready for I/O it should call fuse_notify_poll (possibly asynchronously) with the specified ph; this will clear all pending polls. The callee is responsible for destroying ph with fuse_pollhandle_destroy() when ph is no longer needed.
```

Init Function

The initialization function accepts a `fuse_conn_info` structure, which can be used to investigate and control the system's capabilities. The components of this structure are:

```
proto_major and proto_minor
    Major and minor versions of the FUSE protocol (read-only).

async_read
    On entry, this is nonzero if asynchronous reads are supported. The initialization function can modify this as desired. Note that this field is duplicated by the FUSE_CAP_ASYNC_READ flag; asynchronous reads are controlled by the logical OR of the field and the flag. (Yes, this is a silly hangover from the past.)

max_write
    The maximum size of a write buffer. This can be modified by the init function. If it is set to less than 4096, it is increased to that value.

max_readahead
    The maximum readahead size. This can be modified by the init function.

capabilities
    The capabilities supported by the FUSE kernel module, encoded as bit flags (read-only).

want
    The capabilities desired by the FUSE client, encoded as bit flags.
```

The capabilities that can be requested are:

```
FUSE_CAP_ASYNC_READ
    Use ASYNCREAD reads (default). To disable this option, the client must clear both this capability (in the want flags) and the async_read field. If synchronous reads are chosen, Fuse will wait for reads to complete before issuing any other requests.

FUSE_CAP_POSIX_LOCKS
    Set if the client supports "remote" locking via the lock call.

FUSE_CAP_ATOMIC_O_TRUNC
    Set if the filesystem supports the O_TRUNC open flag.

FUSE_CAP_EXPORT_SUPPORT
    Set if the client handles lookups of "." and ".." itself. Otherwise, FUSE traps these and handles them.

FUSE_CAP_BIG_WRITES
    Set if the filesystem can handle writes larger than 4 KB.

FUSE_CAP_DONT_MASK
    Set to prevent the umask from being applied to files on create operations. (Note: as far as I can tell from examining the code, this flag isn't actually implemented.)
```

Readdir Function

The `readdir` function is somewhat like `read`, in that it starts at a given offset and returns results in a caller-supplied buffer. However, the offset is not a byte offset, and the results are a series of `struct dirent`s rather than being uninterpreted bytes. To make life easier, FUSE provides a "filler" function that will help you put things into the buffer.

The general plan for a complete and correct `readdir` is:

- Find the first directory entry following the given `offset` (see below).
- Optionally, create a `struct stat` that describes the file as the `getattr` (but FUSE only looks at `st_ino` and the file-type bits of `st_mode`).
- Call the `filler` function with arguments of `buf`, the null-terminated filename, the address of your `struct stat` (or NULL if you have none), and the offset of the *next* directory entry. The meaning of "offset" is up to you (see below).
- If `filler` returns nonzero, or if there are no more files, return 0.
- Find the next file in the directory.
- Go back to step 2.

From FUSE's point of view, the `offset` is an uninterpreted `off_t` (i.e., an unsigned integer). You provide an offset when you call `filler`, and it's possible that such an offset might come back to you as an argument later. Typically, it's simply the byte offset (within your directory layout) of the directory entry, but it's really up to you; if you want it to be an index into an array, that's fine.

An important note: you *must* support the `offset` argument. The way Fuse detects the end of the directory is by calling `readdir` with a nonzero `offset` and getting no results back. So if you don't pay attention to the `offset` that is given to you, you'll never terminate. Also, you should never provide a zero `offset`, since Fuse will give that back to you and you'll then start over at the front of the directory!

It's also important to note that `readdir` can return errors in a number of instances; in particular it can return -EBADF if the file handle is invalid, or -ENOENT if you use the `path` argument and the path doesn't exist.

Here's some sample pseudocode:

```
for (i = offset; i < max_dir_entries; i++)
    if (entry[i].valid)
        if (return buf, entry[i].name, NULL, i + 1)
            return 0;
return 0;
```

The above code makes several assumptions and (in general) won't work well with a multi-block directory, but it should give you an outline of how things should work.

Lock function

The [lock](#) function is somewhat complex. The `cmd` will be one of `F_GETLK`, `F_SETLK`, or `F_SETLKW`. The fields in `locks` are defined in the `fcntl(2)` manual page; the `l_whoense` field in that structure will always be `SEEK_SET`.

For checking lock ownership, the `fi->owner` argument must be used.

Contrary to what some other documentation states, the FUSE library does not appear to do anything special to help you out with locking. If you want locking to work, you will need to implement the lock function. (Persons who have more knowledge of how FUSE locking works are encouraged to contact me on this topic, since the existing documentation appears to be inaccurate.)

The Rest of a FUSE Client

Once you've written your operations, you need some boilerplate. As mentioned above, all of your functions should be named with a sensible prefix; here I use "prefix" to represent that. Create a `fuse_operations` struct that lists the functions you implemented (for any unimplemented ones, you could simply delete the relevant lines—but see [Gotchas](#) below for why this is a bad idea):

```
static struct fuse_operations readwrite_oper = {
    .init      = prefix_init,
    .destroy   = prefix_destroy,
    .getattr   = prefix_getattr,
    .fgetattr  = prefix_fgetattr,
    .access    = prefix_access,
    .readlink  = prefix_readlink,
    .readdir   = prefix_readdir,
    .mknod     = prefix_mknod,
    .mkdir     = prefix_mkdir,
    .symlink   = prefix_symlink,
    .unlink    = prefix_unlink,
    .rmdir     = prefix_rmdir,
    .rename    = prefix_rename,
    .link      = prefix_link,
    .chmod     = prefix_chmod,
    .chown     = prefix_chown,
    .truncate  = prefix_truncate,
    .ftruncate = prefix_ftruncate,
    .utimens   = prefix_utimens,
    .create    = prefix_create,
    .open      = prefix_open,
    .read      = prefix_read,
    .write     = prefix_write,
    .stats     = prefix_stats,
    .release   = prefix_release,
    .opendir   = prefix_opendir,
    .releasedir = prefix_releasedir,
    .fsync     = prefix_fsync,
    .flush     = prefix_flush,
    .fsyncdir  = prefix_fsyncdir,
    .lock      = prefix_lock,
    .bmap      = prefix_bmap,
    .ioctl     = prefix_ioctl,
    .poll      = prefix_poll,
#ifdef HAVE_SETXATTR
    .setxattr  = prefix_setxattr,
    .getxattr  = prefix_getxattr,
    .listxattr = prefix_listxattr,
    .removexattr = prefix_removexattr,
#endif
    .flag_nullpath_ok = 0,           /* See below */
};
```

Set `flag_nullpath_ok` nonzero if your code can accept a NULL `path` argument (because it gets file information from `fi->fh`) for the following operations: [fgetattr](#), [flush](#), [fsync](#), [fsyncdir](#), [ftruncate](#), [lock](#), [read](#), [readdir](#), [release](#), [releasedir](#), and [write](#). This will allow FUSE to run more efficiently.

Finally, since your client is actually an executable program, you need a `main`:

```
int main(int argc, char *argv[])
{
    umask(0);
    return fuse_main(argc, &prefix_oper, NULL);
}
```

Compiling Your Program

You can do your development on any machine you choose that supports FUSE. Mac users can try [OSXFUSE](#); Linux users should be able to find FUSE as part of their distribution.

Compiling a FUSE program requires a slightly complicated command:

```
gcc -g -Og `pkg-config fuse --cflags --libs` my_hello.c -o my_hello
```

A better approach, of course, is to use `make`. This truly minimal [Makefile](#) will let you type "make foo" for any `foo.c`. You are encouraged to use it and extend it to be more sensible.

Running & Testing

To run a FUSE program, you'll need two windows and an empty scratch directory. You'll run your filesystem under a debugger in window #1; window #2 will be used for testing. The scratch directory is needed because you must have an empty directory on which to mount your shiny new filesystem.

The simplest (and incorrect, for our purposes) way to run a FUSE program is to make a scratch directory and then pass that as an argument to the program. For example, if you're running the "hello, world" filesystem (`hello.c`):

```
$ mkdir testdir
$ ./hello testdir
$ ls testdir
hello
$ cat testdir/hello
hello, world
$ fusermount -u testdir      # Or on Mac OS, umount testdir
$ rmdir testdir
```

When you run your program this way, it automatically goes into the background and starts serving up your filesystem. After you finish testing, the `fusermount` command unmounts your filesystem and kills the background program.

As a practical matter, it's easier to leave `testdir` lying around rather than making it and removing it every time. (Most production systems have a number of empty directories hanging out just in case you want to mount on top of them—often, either `/mnt` or something inside `/mnt`.)

Of course, it's unlikely that your program will work perfectly the first time, so it's better to run it under the debugger. To do that, you'll need two windows. In window #1, do:

```
$ mkdir testdir      # if necessary
$ gdb hello
(gdb) help
(gdb) [set breakpoints, etc.]
(gdb) run -s -d testdir
```

The `-s` switch means "single-threaded", which makes `gdb` behave in a much friendlier fashion. The `-d` switch means "debug"; in addition to printing helpful debugging output, it keeps the program in the foreground so `gdb` won't lose track of it.

Now, in window #2 you can do:

```
$ ls testdir
$ fusermount -u testdir      # Other trial commands
$ fusermount -u testdir      # Or on Mac OS, umount testdir
```

IMPORTANT: You need to do the `fusermount` even if your program crashes or you abort it. Otherwise you'll get a confusing "Transport endpoint not connected" message the next time you try to mount the test system.

If you have used `gdb` to set breakpoints, then when you do "ls testdir", your window may seem to hang. That's OK; just go over to the `gdb` window and step through your code. When it returns a result, you'll find `testdir` will come alive again.

Other Options

Your new FUSE client has a lot of options. The simplest invocation just specifies a mount point. For example, if your client is named `fuse_client` and you're mounting on `~/foo`, use:

```
./fuse_client ~/foo
```

There are tons of switches available; use `./fuse_client -h` to see them all. The important ones are:

```
-d
    Enable debugging output (implies -f).

-f
    Run in foreground; this is useful if you're running under a debugger. WARNING: When -f is given, Fuse's working directory is the directory you were in when you started it. Without -f, Fuse changes directories to "/". This will screw you up if you use relative pathnames (see "Gotchas" below).

-s
    Run single-threaded instead of multi-threaded. This makes debugging vastly easier, since gdb doesn't handle multiple threads all that well. It also protects you from all sorts of race conditions. Unless you're trying to write a production filesystem and you're a parallelism expert, I recommend that you always use this switch.

-o [no]relinks
    Transform absolute symlinks to relative (or don't, if noellinks is given).
```

Contrary to what the help implies, switches can be specified before the mount point, in standard Unix fashion.

Gotchas

There are several common problems that plague programmers new to Fuse. This is a partial list:

Multithreading

By default, Fuse is multithreaded. That's handy for production filesystems, because it lets client A (or file access A) proceed even if client B is hung up. But multithreading introduces the possibility of race conditions, and makes debugging harder. Always run with the `-s` switch to avoid this problem.

```
getattr
    Fuse calls getattr like crazy. Implement it first, or nothing will work.

Truncate
    Unless you're writing a read-only filesystem, you need to implement the truncate system call to make writes work correctly.
```

Working directory

When it starts, Fuse changes its working directory to `"/`". That will probably break any code that uses relative pathnames, for example while creating a disk file to hold your filesystem. To make matters worse, the `chdir` is suppressed when you run with the `-f` switch, so your code might appear to work fine under the debugger. So if your code creates a scratch file named simply "scratch", with `-f` you'll wind up with "scratch" but without `-f` it'll try to make "scratch", which you don't have permission to do. To avoid the problem, either (a) use absolute pathnames for every file name hardwired into your code (e.g., `~/home/me/fuse/scratch/`), or (b) record your current working directory by calling `get_current_dir_name` before you invoke `fuse_main`, and then convert relative pathnames into corresponding absolute ones. Obviously, (b) is the preferred approach.

Printf

Your `printf`/`fprintf` debugging code will only work if you run with the `-f` switch. Otherwise, Fuse disconnects `stderr` and `stderr`, and all your output will just disappear. So if you use `printf` for debugging, be sure to run with `-f` or `-d`. (Or open a log file with an absolute pathname and use `fprintf` to print to it.)

Unimplemented functions

It is very tempting to just leave functions undefined if your filesystem doesn't need them, or if you just haven't gotten around to writing them yet. Don't! If a function isn't listed in your `fuse_operations` struct, Fuse will silently generate a failure when it is called, and you'll never find out that you need to write it. Instead, write every unimplemented function as a stub that prints a message to `stderr` and returns an error code (ENOYSYS is preferred). When you see the message, you'll know what extra functions you need to write.